

Analyze A/B Test Results

- [Introduction](#)
- [Part I - Probability](#)
- [Part II - A/B Test](#)
- [Part III - Regression](#)
- [Final Check](#)
- [Submission](#)

Introduction

A/B tests are very commonly performed by data analysts and data scientists. For this project, we will be working to understand the results of an A/B test run by an e-commerce website or a company. As a data analyst, my goal for this project is to help the company decide if they should:

- Implement the new webpage,
- Keep the old webpage, or
- Perhaps run the experiment longer to make their decision.

Part I - Probability

To get started, let's import our libraries.

```
#importing libraries
import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline
#We are setting the seed to assure you get the same answers on quizzes as we set up
random.seed(42)
```

1.1

a. Reading the dataset from the `ab_data.csv` file and taking a look at at the top rows:

```
df=pd.read_csv("ab_data.csv")
df.tail()
```

```

        user_id          timestamp    group landing_page
converted
294473    751197  2017-01-03 22:28:38.630509    control    old_page
0
294474    945152  2017-01-12 00:51:57.078372    control    old_page
0
294475    734608  2017-01-22 11:45:03.439544    control    old_page
0
294476    697314  2017-01-15 01:20:28.957438    control    old_page
0
294477    715931  2017-01-16 12:40:24.467417  treatment    new_page
0

```

```
df.columns
```

```
Index(['user_id', 'timestamp', 'group', 'landing_page', 'converted'],
      dtype='object')
```

```
df.describe()
```

```

        user_id    converted
count  294478.000000  294478.000000
mean    787974.124733    0.119659
std      91210.823776    0.324563
min      630000.000000    0.000000
25%      709032.250000    0.000000
50%      787933.500000    0.000000
75%      866911.750000    0.000000
max      945999.000000    1.000000

```

b. Identifying the number of rows in the dataset

```
df.shape
```

```
(294478, 5)
```

c. The number of unique users in the dataset.

```
df.user_id.nunique()
```

```
290584
```

d. The proportion of users converted.

```
df.converted.mean()
```

```
0.11965919355605512
```

```
#total # users converted
```

```
users_converted=df[df['converted']==True].count()
```

```
#find the mean
```

```
users_converted.mean()
```

35237.0

```
#user not converted
```

```
users_not_converted=df[df['converted']==False].count()  
users_not_converted
```

```
user_id      259241  
timestamp    259241  
group        259241  
landing_page  259241  
converted    259241  
dtype: int64
```

e. The number of times when the "group" is treatment but "landing_page" is not a new_page.

```
result_1 = len(df.query('group!="treatment" and  
landing_page=="new_page"))  
result_2 = len(df.query('group!="control" and  
landing_page=="old_page"))  
result_1
```

1928

f. Do any of the rows have missing values?

```
df.isna().sum()
```

```
user_id      0  
timestamp    0  
group        0  
landing_page  0  
converted    0  
dtype: int64
```

1.2

a. Storing dataframe in df2

```
# Remove the inaccurate rows, and store the result in a new dataframe  
df2
```

```
df2 = df[((df['group']=='treatment') &  
(df['landing_page']=='new_page')) | ((df['group']=='control') &  
(df['landing_page']=='old_page'))]
```

```
# Double Checking if all of the incorrect rows were removed from df2  
df2[((df2['group'] == 'treatment') == (df2['landing_page'] ==  
'new_page')) == False].shape[0]
```

0

1.3

a. How many unique **user_ids** are in **df2**?

```
df2['user_id'].nunique()
```

```
290584
```

b. There is one **user_id** repeated in **df2**. What is it?

```
df2[df2['user_id'].duplicated()][ 'user_id' ]
```

```
2893      773192
```

```
Name: user_id, dtype: int64
```

c. Display the rows for the duplicate **user_id**?

```
df2[df2['user_id'].duplicated()]
```

	user_id	timestamp	group	landing_page
converted				
2893	773192	2017-01-14 02:55:59.590927	treatment	new_page
0				

d. Remove **one** of the rows with a duplicate **user_id**, from the **df2** dataframe.

```
# Remove one of the rows with a duplicate user_id..
```

```
df2 = df2.drop_duplicates('user_id');
```

```
df2[df2['user_id'].duplicated()][ 'user_id' ]
```

```
Series([], Name: user_id, dtype: int64)
```

1.4

a. What is the probability of an individual converting regardless of the page they receive?

```
df2['converted'].mean()
```

```
0.11959708724499628
```

b. Given that an individual was in the control group, what is the probability they converted?

```
df2.query('group == "control"').converted.mean()
```

```
0.1203863045004612
```

c. Given that an individual was in the treatment group, what is the probability they converted?

```
df2.query('group == "treatment"').converted.mean()
```

```
0.11880806551510564
```

d. What is the probability that an individual received the new page?

```
len(df2[df2['landing_page'] == 'new_page'])/len(df2)
```

0.5000619442226688

e. Consider the results above, we cannot really tell the difference between the groups group control had more conversion compared to experiment group yet there's no concrete evidence to prove it.

Part II - A/B Test

2.1

For now, consider you need to make the decision just based on all the data provided.

Recall that we just calculated that the "converted" probability (or rate) for the old page is *slightly* higher than that of the new page (ToDo 1.4.c).

Hypothesis:

2.2 - Null Hypothesis H_0 Testing

Under the null hypothesis H_0 , assume that p_{new} and p_{old} are equal. Furthermore, assume that p_{new} and p_{old} both are equal to the **converted** success rate in the df2 data regardless of the page. So, our assumption is: $p_{new} = p_{old} = p_{population}$

In this section, we will:

- Simulate (bootstrap) sample data set for both groups, and compute the "converted" probability p for those samples.
- Use a sample size for each group equal to the ones in the df2 data.
- Compute the difference in the "converted" probability for the two samples above.
- Perform the sampling distribution for the "difference in the converted probability" between the two simulated-samples over 10,000 iterations; and calculate an estimate.

a. What is the **conversion rate** for p_{new} under the null hypothesis?

```
p_new = df2.converted.mean()
```

p_new

0.11959708724499628

b. What is the **conversion rate** for p_{old} under the null hypothesis?

```
p_old = df2.converted.mean()  
p_old
```

```
0.11959708724499628
```

c. What is n_{new} , the number of individuals in the treatment group?

```
#n_new = df2.landing_page.value_counts()[0]  
#n_new  
new_df = df2.query('landing_page == "new_page"')  
n_new = new_df.shape[0]  
n_new
```

```
145310
```

d. What is n_{old} , the number of individuals in the control group?

```
n_old = df2.landing_page.value_counts()[1]  
n_old  
  
older_df = df2.query('landing_page == "old_page"')  
n_old = older_df.shape[0]  
n_old
```

```
145274
```

e. **Simulate Sample for the treatment Group** Simulate n_{new} transactions with a conversion rate of p_{new} under the null hypothesis.

```
# Simulate a Sample for the treatment Group  
new_page_converted = np.random.binomial(n_new, p_new)  
new_page_converted
```

```
17367
```

f. **Simulate Sample for the control Group** Simulate n_{old} transactions with a conversion rate of p_{old} under the null hypothesis. Store these n_{old} 1's and 0's in the old_page_converted numpy array.

```
old_page_converted = np.random.binomial(n_old, p_old)
```

g. Find the difference in the "converted" probability $\hat{p} - p'_{old}\hat{p}$ for your simulated samples from the parts (e) and (f) above.

```
p_diff = (new_page_converted/n_new) - (old_page_converted/n_old)  
p_diff  
p_diff = (new_page_converted/n_new) - (old_page_converted/n_old)  
p_diff
```

```
-0.00031872605020111244
```

h. Sampling distribution Re-create new_page_converted and old_page_converted and find the $\hat{p} - p'_{old}$ value 10,000 times using the same simulation process you used in parts (a) through (g) above.


Store all $\hat{p} - p'_{old}$ values in a NumPy array called p_diffs.

```
p_diffs = []
```

```
for _ in range(10000):
    new_converted_simulation = np.random.binomial(n_new,p_new)/n_new
    old_converted_simulation = np.random.binomial(n_old,p_old)/n_old
    diff = new_converted_simulation - old_converted_simulation
    p_diffs.append(diff)
```

i. Histogram

```
plt.hist(p_diffs);
```



```
import statsmodels.api as sm
convert_old = sum(df2.query("group == 'control'")['converted'])
convert_new = sum(df2.query("group == 'treatment'")['converted'])
n_old = len(df2.query("group == 'control'"))
n_new = len(df2.query("group == 'treatment'"))

z_score, p_value = sm.stats.proportions_ztest([convert_old,
convert_new], [n_old, n_new], alternative='smaller')
z_score, p_value
(1.3109241984234394, 0.9050583127590245)
```

```
from scipy.stats import norm
norm.cdf(z_score)
norm.ppf(1-(0.05))
```

1.6448536269514722


j. What proportion of the p_diffs are greater than the actual difference observed in ab_data.csv?

```
old_mean = df.query('group == "control"').converted.mean()
new_mean = df.query('group == "treatment"').converted.mean()
a_diff = new_mean - old_mean
p_diffs = np.array(p_diffs)
(p_diffs > a_diff).mean()
```

0.8893

```
plt.hist(p_diffs)
plt.xlabel('p_diffs')
plt.ylabel('Frequency')
plt.title('Simulated Difference between new_page & old_page ')
plt.axvline(x=a_diff, color='r', label="Observed difference")
```

<matplotlib.lines.Line2D at 0x2948fc8aa60>



k. In words, explain what you just computed in part j.. What is this value called in scientific studies? What does this value mean in terms of whether or not there is a difference between the new and old pages?

The value above represents the p-value of observing the statistic given the Null is true. As the p-value is large enough, we would fail to reject the Null hypothesis and keep the old page.

We could also use a built-in to achieve similar results. Though using the built-in might be easier to code, the above portions are a walkthrough of the ideas that are critical to correctly thinking about statistical significance. Fill in the below to calculate the number of conversions for each page, as well as the number of individuals who received each page. Let `n_old` and `n_new` refer to the number of rows associated with the old page and new pages, respectively.

```
import statsmodels.api as sm

convert_old = sum(df2.query("group == 'control'")['converted'])
convert_new = sum(df2.query("group == 'treatment'")['converted'])
n_old = len(df2.query("group == 'control'"))
n_new = len(df2.query("group == 'treatment'"))

z_score, p_value = sm.stats.proportions_ztest([convert_old,
convert_new], [n_old, n_new], alternative='smaller')

z_score, p_value

(1.3109241984234394, 0.9050583127590245)
```

n. What do the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? Do they agree with the findings in parts j. and k.?

```
norm.cdf(z_score)

0.9050583127590245

norm.ppf(1 - (0.05))

1.6448536269514722
```

Observation

Part III - A regression approach

3.1

Since each row is either a conversion or no conversion, what type of regression should you be performing in this case?

b. The goal is to use **statsmodels** library to fit the regression model you specified in part **a.** above to see if there is a significant difference in conversion based on the page-type a

customer receives. However, you first need to create the following two columns in the df2 dataframe:

1. `intercept` - It should be 1 in the entire column.
2. `ab_page` - It's a dummy variable column, having a value 1 when an individual receives the **treatment**, otherwise 0.

```
df2['intercept']=1
df2[['control', 'treatment']] = pd.get_dummies(df2['group'])
```

c. Use **statsmodels** to instantiate your regression model on the two columns you created in part (b). above, then fit the model to predict whether or not an individual converts.

```
import statsmodels.api as sm
logit = sm.Logit(df2['converted'],df2[['intercept','treatment']])
```

d. Provide the summary of your model below, and use it as necessary to answer the following questions.

```
results = logit.fit()
results.summary()
```

```
Optimization terminated successfully.
      Current function value: 0.366118
      Iterations 6
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

Logit Regression Results

```
=====
=====
Dep. Variable:          converted    No. Observations:
290584
Model:                  Logit        Df Residuals:
290582
Method:                 MLE          Df Model:
1
Date:                   Sat, 18 Feb 2023    Pseudo R-squ.:
8.077e-06
Time:                   10:14:19    Log-Likelihood:    -
1.0639e+05
converged:              True          LL-Null:            -
1.0639e+05
Covariance Type:        nonrobust    LLR p-value:
0.1899
=====
=====
               coef      std err          z      P>|z|      [0.025
0.975]
-----
-----
```

```

intercept      -1.9888      0.008    -246.669      0.000      -2.005
-1.973
treatment      -0.0150      0.011     -1.311      0.190      -0.037
0.007
=====
=====
" " "

```

There difference observed in this project is that in part II, the focus was on one side test, whereas in part III focus was on two side test

f. Now, you are considering other things that might influence whether or not an individual converts. Discuss why it is a good idea to consider other factors to add into your regression model. Are there any disadvantages to adding additional terms into your regression model?

We could add variables to determine whether they will have influence or not. However, adding to many features into the regression model can result in over-fitting.

We could also investigate device if they had influence on conversion

g. Adding countries Adding countries to observe if the results will be impacted by where the user lives. Now along with testing if the conversion rate changes for different pages, also add an effect based on which country a user lives in.

```

# Read the countries.csv
countries_df = pd.read_csv("countries.csv")
countries_df =
countries_df.set_index('user_id').join(df2.set_index('user_id'),
how='inner')
countries_df.head()

```

user_id	country	timestamp	group	landing_page \
834778	UK	2017-01-14 23:08:43.304998	control	old_page
928468	US	2017-01-23 14:44:16.387854	treatment	new_page
822059	UK	2017-01-16 14:04:14.719771	treatment	new_page
711597	UK	2017-01-22 03:14:24.763511	control	old_page
710616	UK	2017-01-16 13:14:44.000513	treatment	new_page

user_id	converted	intercept	control	treatment
834778	0	1	1	0
928468	0	1	0	1
822059	1	1	0	1
711597	0	1	1	0
710616	0	1	0	1

```

# Join with the df2 dataframe
countries_df['country'].value_counts()

```

```
US      203619
UK       72466
CA       14499
Name: country, dtype: int64
```

```
# Create the necessary dummy variables
```

```
countries_df[['CA', 'US']] = pd.get_dummies(countries_df['country'])
[['CA', 'US']]
```

h. Fit your model and obtain the results Though you have now looked at the individual factors of country and page on conversion, we would now like to look at an interaction between page and country to see if are there significant effects on conversion.

```
# Fit your model, and summarize the results
```

```
log_mod = sm.Logit(countries_df['converted'], countries_df[['CA',
'US']])
results = log_mod.fit()
results.summary()
```

```
Optimization terminated successfully.
      Current function value: 0.447174
      Iterations 6
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

Logit Regression Results

```
=====
=====
```

Dep. Variable:	converted	No. Observations:
290584		
Model:	Logit	Df Residuals:
290582		
Method:	MLE	Df Model:
1		
Date:	Sat, 18 Feb 2023	Pseudo R-squ.:
-0.2214		
Time:	10:14:23	Log-Likelihood:
1.2994e+05		-
converged:	True	LL-Null:
1.0639e+05		-
Covariance Type:	nonrobust	LLR p-value:
1.000		

```
=====
=====
```

	coef	std err	z	P> z	[0.025
--	------	---------	---	------	--------

```
0.975]
```

```
-----
-----
```

CA	-2.0375	0.026	-78.364	0.000	-2.088
-1.987					

US	-1.9967	0.007	-292.314	0.000	-2.010
-1.983					

```
=====
=====
"""
```

Based on the analysis, we conclude neither new page or old page will lead to higher conversion and for this reason we say that we fail to reject the null hypothesis thus we advise the company to keep the old page or run the test for sometimes.

```
from subprocess import call
call(['python', '-m', 'nbconvert',
'Analyze_ab_test_results_notebook.ipynb'])
```

1